

Adapting a Latent Audio Diffusion Model to Historical Guqin Recordings: A Listening-Driven Case Study

Huanchen Cai*

October 2026

Abstract

We report a small-data case study in adapting a pretrained latent audio diffusion model to the guqin, the seven-string Chinese zither, with the long-term goal of an “AI radio” that plays guqin-style music without end. From a library of historical recordings we curate 412 solo performances (42.5 h, 61 performers) and split them by composition. We fine-tune a rank-16 DoRA adapter on Stable Audio 3 Medium using its continuous SAME-L latents, continuation-oriented masking, captions combining researched notes on each piece with mood tags and an automatically estimated pentatonic mode, and random-length crops, and we stop training when held-out validation loss stops improving. Seven blind listening studies by one expert listener guide every decision. The final adapter (pass 14 of 17) received the highest mean rating for continuing unseen pieces (3.9/5 versus 3.4 for the best earlier adapter and 2.8 for an early checkpoint of the same run) and 4.6/5 for generating from free-written scene descriptions. Equally informative are the failures: a from-scratch autoregressive model over the same latents produced no recognisable timbre; a signal-level “friction noise” measure correlated with the listener’s “sawing” complaints in the wrong direction; and a pentatonic-fit measure tracked ratings overall ($\rho = 0.52$) but barely within a group of candidates. Chaining continuations for long playback exposed two mechanical problems, a silent tail on every generated clip and a loudness feedback loop, both of which have simple fixes. With a single listener and at most ten clips per condition, none of the paired differences is statistically significant; we present the results as an exploratory record of what helped, what did not, and why.

1 Introduction

Text-to-music and audio continuation models are trained mostly on Western popular and library music. The guqin is a poor fit for that prior: its repertoire is largely pentatonic, its expressive vocabulary lives in left-hand slides and vibrato (吟猱绰注) whose passing semitones are idiomatic only in specific places, and the best-known performances survive mainly as mid-twentieth-century recordings with tape hiss and limited bandwidth. The amount of usable solo audio is also small—tens of hours, not thousands.

Our practical goal is an “AI radio”: a system that keeps producing guqin-like music for hours, either continuing from a recorded opening or following a short textual description of a scene. This report documents how we got from a general-purpose model to an adapter that a guqin-literate listener rates as convincing, and records the dead ends along the way, because in a domain this small the negative results are a large part of what was learned.

*[TODO: affiliation and contact e-mail]

Our contributions are:

- A curation and evaluation protocol for a small historical corpus: expert screening of every recording, composition-level splits so that test pieces are never heard in another performer’s version, and blind, shuffled listening with free-text comments (§3, §5).
- A fine-tuning recipe for Stable Audio 3 Medium [6] on 42 h of audio: DoRA [10] on the base model with inference on the distilled model, masking weighted towards prefix continuation, captions that mix researched notes on each composition with mood tags and an estimated pentatonic mode, random-length crops, and early stopping on held-out loss (§4).
- Seven listening studies and the decisions they drove, including clear negative results for a from-scratch latent language model and for two automatic quality measures (§6).
- Two practical failure modes of chained continuation for long playback, with fixes (§7).

2 Related work

Latent audio diffusion. Stable Audio [3] generates variable-length audio by diffusion in the latent space of an autoencoder, conditioned on text and timing; Stable Audio Open [5] and Stable Audio 3 [6] extend this with open weights, inpainting and continuation. Long-form generation with latent diffusion [4] shows that the training window length matters for musical structure. We build directly on Stable Audio 3 Medium and its continuous SAME-L autoencoder.

Autoregressive music models. Token-based language models such as MusicGen [2] and streaming systems such as Magenta RealTime [12] generate indefinitely by construction. Continuous-latent autoregression with a diffusion or flow head [9] avoids quantisation, and noise augmentation of the context reduces error accumulation [14]. We test a small model of this kind trained from scratch (§6.1).

Parameter-efficient adaptation and small data. LoRA [7] and DoRA [10] adapt large models with few trainable parameters. For data-constrained training, repeating data for a few epochs is nearly as useful as fresh data [13], which motivates our multi-pass schedule with early stopping.

Guqin and evaluation. Prior guqin MIR work analyses mode and playing technique from historical recordings [8]. For generated music, replication of training data is a central concern [1]; we keep evaluation pieces out of training at the composition level.

3 Data

Source and screening. The source is a private library of 638 historical guqin recordings (44.1 kHz, 16-bit stereo CD transfers). Every candidate was screened by the author, a guqin-literate listener, through a whole-track listening page; recordings were excluded for audible noise floors, ensembles or non-guqin instruments, and voice. The result is **412 solo recordings by 61 performers, 42.5 h**, with 108 distinct titles after normalising numbered variants (a few titles are alternative names of the same piece). Performer coverage is uneven: the five most represented performers account for about 40% of the duration.

Splits. We split by composition rather than by performer, so that no version of a test composition is heard in training: 329 recordings for training, 42 for validation and 41 for test. Each recording is cut into aligned windows of the model’s native length (59.9 s, 2,641,920 samples) that together cover the whole performance, giving 2,764 windows (2,217/276/271); nine windows of pure digital silence are dropped. The final run trains on training and test recordings together (370 recordings) and keeps the 42 validation recordings out for early stopping and listening.

Captions. For each composition we compiled a short narrative from published notes on the piece (e.g. *Liushui*, Boya and his true listener Ziqi; *Guangling San*, Nie Zheng’s revenge and Ji Kang’s last performance), assigned two to five tags from a closed vocabulary of 37 moods and scenes (e.g. *calm*, *lonely*, *vast*, *reclusive*, *autumn*, *river*, *frontier desert*), and recorded our confidence: 55 notes are well sourced, 31 partly, and 22 inferred from the title. A pentatonic mode (the *gong* pitch and its five notes) is estimated automatically for every window from a tuning-corrected pitch-class histogram. Captions are written in English with the Chinese text appended and never exceed 160 tokens of the 256-token encoder limit.

4 Method

Model. Stable Audio 3 Medium [6] is a diffusion transformer over SAME-L latents (256 dimensions at about 10.8 frames per second) conditioned on T5Gemma text embeddings and on an inpainting mask. Following the released recipe, we train a rank-16 DoRA adapter (21.6 M trainable parameters) on the *base* model and apply it to the step-distilled *Medium* model at inference (8 sampling steps). All latents are pre-encoded once.

Objective and masking. Training uses the model’s rectified-flow objective [11] with its inpainting conditioning. We change the masking mixture from the default 10/80/10 (random spans / fully masked / prefix continuation) to 10/30/60, so that most training examples resemble our main use, continuing a given opening.

Text. Each window receives one caption style by a stable hash: 40% tags + mode, 40% narrative + tags + mode, 20% narrative only; the trainer additionally drops 10% of captions for classifier-free use.

Crops and schedule. Each sample keeps a random span of 20–60 s at a random start inside its window, padded with the silence latent and with the timing condition set to the kept length. One training pass is 2,500 steps at batch size 1 (about one epoch), learning rate 10^{-4} , AdamW. After each pass we compute the model’s validation loss at five fixed noise levels on all 276 validation windows and stop when three consecutive passes fail to improve on the best. Checkpoints every 500 steps make the run resumable after interruptions.

Inference. Continuation keeps a 10 s opening and generates the next 25 s; text-only generation produces 35 s from a description. For long playback we chain continuations (§7).

5 Evaluation protocol

All listening was done by the author, who has extensive experience with the guqin repertoire, on local web pages in a fixed format: candidates for the same opening and seed are shuffled and

Study	Condition	n	Mean
R1 masking (329 vs 412)	default mask / continuation mask / 412, cont. mask	8	2.88 / 2.62 / 3.38
R2 inference settings	strength 1.0 / 0.7 / LoRA early steps only	9–10	3.56 / 3.60 / 3.67
	base model, 50 steps / earlier 222-clip adapter	10	2.70 / 3.10
R3 captions (continuation)	no text in training / no prompt / matching / opposite tags	9	3.89 / 3.56 / 4.00 / 4.00
R3 captions (text only)	adapter without / with caption training	5	2.00 / 3.80
R4 latent LM	from-scratch latent LM / SA3 adapter	11 / 6	1.00 / 2.00
R5 longer training	1 pass / 3 passes + tags / 3 passes + tags + mode	8	3.75 / 3.75 / 4.12
	pass 2 / adapter C (R5)	7 / 8	3.00 / 3.50
R7 final	pass 14 / pass 2 / adapter C	8	3.88 / 2.75 / 3.38
R7 final (scenes)	pass 14 / pass 2	5	4.60 / 3.60

Table 1: Overview of the listening studies (scores 1–5, single expert listener). In R1 the 412-recording adapter had seen the validation compositions; all other comparisons use openings unseen by every compared model, except R2 (see §8).

unlabeled; each clip gets an overall score from 1 to 5, optional issue tags (e.g. *too many semitones, sawing/metallic noise, not guqin-like, no melody, abrupt seam*) and free-text comments; the mapping is revealed only after rating. Openings come from compositions the compared models did not train on, with each study using different performers and pieces where possible. For paired comparisons we report exact two-sided sign-flip permutation tests over openings and percentile bootstrap intervals; with eight or fewer pairs these tests have little power and we treat them as descriptive.

6 Results

Table 1 summarises the seven studies. Scores are means over clips; n is clips per condition.

Seed and opening dominate. In R1 the two seeds differed by 1.6 points on average (2.17 vs 3.75), more than any model difference; in R2 the mean per opening-and-seed group ranged from 1.0 to 4.75 while inference settings moved the mean by at most 0.1 (except the base model). The same semitone often appeared at the same second in all candidates sharing an opening and seed. This variance motivates generating several takes and selecting, and it is why per-condition means over a handful of clips must be read cautiously.

Inference settings. Lowering adapter strength or restricting the adapter to early denoising steps made no difference. Using the base model with 50 steps—the configuration matching training—was clearly worse (2.70, five scores of 1, comments about extra instruments), confirming that the distilled model is the better renderer.

Text conditioning. Training with mood tags left continuation unchanged (matching and deliberately opposite tags both scored 4.0: the opening dominates), but improved generation from text alone from 2.0 to 3.8 ($p = 0.06$, $n = 5$).




Figure 1: Validation loss per training pass of the final run (276 held-out windows, five fixed noise levels).

Training length. On the 329-piece split, held-out loss decreased monotonically over three passes (0.5387 base, 0.5351, 0.5343, 0.5337) with no sign of overfitting; captions and mode text did not change the loss. In the final run (Figure 1) the loss fell from 0.6020 to a minimum of 0.5950 at pass 14 and rose over the next three passes, triggering early stopping. The pass-14 adapter improved continuation over the pass-2 checkpoint of the same run (3.88 vs 2.75, $p = 0.16$) and over adapter C (3.38, $p = 0.63$), and received no “sawing” tags against three for pass 2. Bootstrap 95% intervals are wide ($[2.75, 4.75]$ for pass 14). For scene descriptions written after training, pass 14 scored 4.6 against 3.6 ($p = 0.31$); one take was singled out by the listener as exceptional.

6.1 Negative result: a from-scratch latent language model

We trained a causal transformer (8 layers, width 512) with a flow-matching MLP head that predicts the next SAME-L frame, with context noise augmentation [14] (72 M parameters, 1.42 M training frames). Validation loss reached its minimum at 10k steps and rose by 8% by 40k steps (Figure 2). Every one of its 11 clips received the lowest score, including endless and unconditioned generation, with comments such as “no timbre at all”; independently sampled frames left the region the decoder renders as guqin. The same openings continued by the SA3 adapter scored 2.0. At this data scale a frame-level autoregressive model over continuous latents is not competitive with adapting a large pretrained generator.

6.2 Negative and weak results: automatic quality measures

Friction-noise measure. The listener’s most frequent complaint was “sawing”: left-hand friction noise that is idiomatic in moderation but grating when overused. A measure of noisy high-band frames correlated with the score in the wrong direction (Spearman $\rho = +0.39$ over 48 clips), and the nine clips tagged as sawing had *lower* noisy-frame fractions than the rest (0.0006 vs 0.0017).

Pentatonic fit. The fraction of tonal energy on the best-fitting pentatonic set tracked scores over 124 rated clips ($\rho = 0.52$, positive in each study) but only weakly among candidates for the same opening ($\rho = 0.31$ after removing group means). Choosing the best-fitting candidate raised the mean from 3.41 to 3.52; discarding the worst raised it to 3.58. Restricting the measure to pitches at note onsets—motivated by the listener’s observation that idiomatic semitones occur in slides, not at




Figure 2: From-scratch latent LM: validation loss rises after 10k steps (dashed).

the pluck—was weaker still ($\rho_{\text{within}} = 0.19$), although generated clips did place more onsets outside the mode than real recordings (33% vs a median of 22%; 39% in clips flagged for semitones). The measure on its own also disagreed with R5, where it predicted that longer training worsened modal fit while the listener preferred the longer-trained model.

6.3 Rater consistency

Fifteen clips from R6 were rated again, unknowingly, in R7. Six scores were identical, twelve within one point, and the mean absolute difference was 0.8. A clip chosen for the showcase in one session received a 1 in the next when heard next to an alternative. Context effects of this size are comparable to the differences between our models and are the main reason we do not claim significance.

7 Long playback by chained continuation

To extend the favourite 35 s clip, we repeatedly fed its last 15 s as an opening and appended 30 s of new material. Two problems appeared.

Silent tails. Every generated clip ends with about 0.5 s of digital silence (-87 dBFS) as the requested duration runs out. When that tail is part of the next opening, the model treats the piece as ended and inserts a further two seconds of silence at the seam. Trimming the source’s last 1.5 s and over-generating each segment by 3 s, then discarding the tail, removed the gaps.

Loudness feedback. Each segment came out slightly louder than its opening, and that segment became the next opening: over six rounds the level rose from -10 to -5.5 dBFS with up to 10% of samples clipped. Feeding every opening at a fixed level (-16 dBFS) and matching each new segment to the level of the 15 s before it kept the 3.5-minute result between -13.8 and -16.7 dBFS per 30 s with no clipping. Of two locked chains, the listener accepted one and rejected the other for a sudden breakdown after 82 s, a reminder that long-range stability is not solved.

8 Limitations and ethics

Evaluation. All ratings come from one listener, who is also the author and designed the systems; the studies are small and were run sequentially as development proceeded, so later studies are not independent of earlier decisions. In R2 the openings came from a performer who, contrary to the original plan, had been included in the 412-recording training set. Multi-listener studies, objective measures validated against them, and checks for replication of training material [1] are needed before stronger claims.

Data and rights. The training recordings are commercially published historical performances and are not redistributed; we do not release the adapter weights. We release code, the listening ratings and the analysis that produces every number in this report [TODO: repository URL]. The base model is used under Stability AI’s community license. Openly licensed guqin recordings we could find amount to under one hour.

Cultural context. The narratives and tags are our condensed readings of traditional notes on each piece and simplify a rich interpretive tradition; 22 of the 108 titles are inferred from the title alone.

9 Conclusion

A pretrained latent diffusion model, a few tens of hours of carefully screened historical recordings, a small adapter and a single attentive listener were enough to produce guqin continuations and scene-driven pieces that the listener rated highly, while training a generator from scratch on the same data failed completely. The strongest effects were not from architecture but from data curation, held-out early stopping, and simple engineering of the generation loop. The next steps are multi-listener evaluation, preference-based fine-tuning on the collected ratings [15], and stable long-form generation for continuous playback.

Acknowledgements

Code for data processing, training orchestration, evaluation pages and analysis was written with the help of AI coding assistants (OpenAI Codex and Anthropic Claude); all experimental decisions and all listening judgements are the author’s. [TODO: other acknowledgements]

References

- [1] Roser Batlle-Roca, Wei-Hsiang Liao, Xavier Serra, Yuki Mitsufuji, et al. Towards assessing data replication in music generation with music similarity metrics on raw audio. *arXiv preprint arXiv:2407.14364*, 2024.
- [2] Jade Copet, Felix Kreuk, Itai Gat, Tal Remez, et al. Simple and controllable music generation. *arXiv preprint arXiv:2306.05284*, 2023.
- [3] Zach Evans, CJ Carr, Josiah Taylor, Scott H. Hawley, and Jordi Pons. Fast timing-conditioned latent audio diffusion. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2402.04825.

- [4] Zach Evans, Julian D. Parker, CJ Carr, Zack Zukowski, Josiah Taylor, and Jordi Pons. Long-form music generation with latent diffusion. *arXiv preprint arXiv:2404.10301*, 2024.
- [5] Zach Evans, Julian D. Parker, CJ Carr, Zack Zukowski, Josiah Taylor, and Jordi Pons. Stable audio open. *arXiv preprint arXiv:2407.14358*, 2024.
- [6] Zach Evans, Julian D. Parker, Matthew Rice, et al. Stable audio 3. *arXiv preprint arXiv:2605.17991*, 2026.
- [7] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- [8] Yu-Fen Huang, Jeng-I Liang, I-Chieh Wei, and Li Su. Joint analysis of mode and playing technique in guqin performance with machine learning. In *Proceedings of the International Society for Music Information Retrieval Conference (ISMIR)*, 2020.
- [9] Tianhong Li, Yonglong Tian, He Li, Mingyang Deng, and Kaiming He. Autoregressive image generation without vector quantization. *arXiv preprint arXiv:2406.11838*, 2024.
- [10] Shih-Yang Liu, Chien-Yi Wang, Hongxu Yin, Pavlo Molchanov, et al. DoRA: Weight-decomposed low-rank adaptation. *arXiv preprint arXiv:2402.09353*, 2024.
- [11] Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. *arXiv preprint arXiv:2209.03003*, 2022.
- [12] Lyria Team, Antoine Caillon, Brian McWilliams, Cassie Tarakajian, et al. Live music models. *arXiv preprint arXiv:2508.04651*, 2025.
- [13] Niklas Muennighoff, Alexander M. Rush, Boaz Barak, Teven Le Scao, et al. Scaling data-constrained language models. *arXiv preprint arXiv:2305.16264*, 2023.
- [14] Marco Pasini, Javier Nistal, Stefan Lattner, and George Fazekas. Continuous autoregressive models with noise augmentation avoid error accumulation. *arXiv preprint arXiv:2411.18447*, 2024.
- [15] Bram Wallace, Meihua Dang, Rafael Rafailov, Linqi Zhou, et al. Diffusion model alignment using direct preference optimization. *arXiv preprint arXiv:2311.12908*, 2023.